

Dataset Operations

DATA IMPORT / EXPORT

h2o.downloadCSV: Download a H2O dataset to a CSV file on local disk.

h2o.exportFile: Export H2O Data Frame to a file.

h2o.importFile: Import a file from the local path and parse it.

h2o.parseRaw: Parse a raw data file.

h2o.uploadFile: Upload a file from the local drive and parse it.

NATIVE R TO H2O COERCION

as.h2o: Convert an R object to an H2O object.

H2O TO NATIVE R COERCION

as.data.frame: Check if an object is a data frame, or coerce it if possible.

DATA GENERATION

h2o.createFrame: Create an H2O data frame, with optional randomization.

h2o.runif: Produce a vector of random uniform numbers.

h2o.interaction: Create interaction terms between categorical features of an H2O Frame.

DATA SAMPLING / SPLITTING

h2o.splitFrame: Split an existing H2O dataset according to user-specified ratios.

MISSING DATA HANDLING

h2o.impute: Impute a column of data using the mean, median, or mode.

h2o.insertMissingValues: Replaces a user-specified fraction of entries in a H2O dataset with missing values.

General Operations

SUBSCRIPTING

Subscripting example to pull pieces from data object.

<pre>x[j] ## column J x[i, j] x[[i]] x\$name</pre>	Selection	<pre>x[i] <- value x[i, j, ...] <- value x[[i]] <- value x\$i <- value</pre>	Value Assignment
--	-----------	--	------------------

SUBSETTING

head, tail: Return the First or Last Part of an Object

CONCATENATION

c: Combine Values into a Vector or List.

h2o.cbind: Take a sequence of H2O datasets and combine them by column.

DATA ATTRIBUTES

colnames: Return column names for a parsed H2O data object.

colnames<-: Retrieve or set the row or column names of a matrix-like object.

names: Get the name of an object.

names<-: Set the name of an object.

dim: Retrieve the dimension of an object.

length: Get the length of vectors (including lists) and factors.

nrow: Return a count of the number of rows in an H2OParsedData object.

ncol: Return a count of the number of columns in an H2OParsedData object.

h2o.anyFactor: Check if an H2O parsed data object has any categorical data columns.

is.factor: Check if a given column contains categorical data.

DATA TYPE COERCION

as.factor: Convert a column from numeric to factor.

as.Date: Converts a column from factor to date.

Methods from Group Generics: Math

MATH (H2O)

abs: Compute the absolute value of x.

sign: Return a vector with the signs of the corresponding elements of x (the sign of a real number is 1, 0, or -1 if the number is positive, zero, or negative, respectively).

sqrt: Computes the principal square root of x, \sqrt{x} .

ceiling: Take a single numeric argument x and return a numeric vector containing the smallest integers not less than the corresponding elements of x.

floor: Take a single numeric argument x and return a numeric vector containing the largest integers not greater than the corresponding elements of x.

trunc: Take a single numeric argument x and return a numeric vector containing the integers formed by truncating the values in x toward 0.

log: Compute logarithms (by default, natural logarithms).

exp: Compute the exponential function.

MATH (GENERIC)

cummax: Display a vector of the cumulative maxima of the elements of the argument.

cummin: Display a vector of the cumulative minima of the elements of the argument.

cumprod: Display a vector of the cumulative products of the elements of the argument.

cumsum: Display a vector of the cumulative sums of the elements of the argument.

log10: Compute common (i.e., base 10) logarithms.

log2: Compute binary (i.e., base 2) logarithms.

log1p: Compute $\log(1+x)$ accurately also for $|x| \ll 1$.

MATH (GENERIC)

acos: Compute the trigonometric arc-cosine.

acosh: Compute the hyperbolic arc-cosine.

asin: Compute the trigonometric arc-sine.

asinh: Compute the hyperbolic arc-sine.

atan: Compute the trigonometric arc-tangent.

atanh: Compute the hyperbolic arc-tangent.

expm1: Compute $\exp(x) - 1$ accurately also for $|x| \ll 1$.

cos: Compute the trigonometric cosine.

cosh: Compute the hyperbolic cosine.

cospi: Compute the trigonometric two-argument arc-cosine.

sin: Compute the trigonometric sine.

sinh: Compute the hyperbolic sine.

sinpi: Compute the trigonometric two-argument arc-sine.

tan: Compute the trigonometric tangent.

tanh: Compute the hyperbolic tangent.

tanpi: Compute the trigonometric two-argument arc-tangent.

gamma: Display the gamma function γx

lgamma: Display the natural logarithm of the absolute value of the gamma function.

digamma: Display the first derivative of the logarithm of the gamma function.

trigamma: Display the second derivative of the logarithm of the gamma function.

MATH2 (H2O)

round: Round the values to the specified number of decimal places. The default is 0.

signif: Round the values to the specified number of significant digits.

Methods from Group Generics: Summary

SUMMARY (H2O)

max: Display the maximum of all the input arguments.

min: Display the minimum of all the input arguments.

range: Display a vector containing the minimum and maximum of all the given arguments.

sum: Calculate the sum of all the values present in its arguments.

SUMMARY (GENERIC)

prod: Display the product of all values present in its arguments.

any: Given a set of logical vectors, determine if at least one of the values is true.

all: Given a set of logical vectors, determine if all of the values are true.

Other Aggregations

NON-GROUP GENERIC SUMMARIES

mean: Generic function for the (trimmed) arithmetic mean.

sd: Calculate the standard deviation of a column of continuous real valued data.

var: Compute the variance of x.

summary: Produce result summaries of the results of various model fitting functions.

quantile: Obtain and display quantiles for H2O parsed data.

ROW / COLUMN AGGREGATION

apply: Apply a function over an H2O parsed data object (an array).

GROUP BY AGGREGATION

h2o.group by: Apply an aggregate function to each group of an H2O dataset.

TABULATION

h2o.table: Use the cross-classifying factors to build a table of counts at each combination of factor levels.

Data Munging

GENERAL COLUMN MANIPULATION

is.na: Display missing elements.

ELEMENT INDEX SELECTION

h2o.which: Display the row numbers for which the condition is true.

CONDITIONAL ELEMENT VALUE SELECTION

h2o.ifelse: Apply conditional statements to numeric vectors in H2O parsed data objects.

NUMERIC COLUMN MANIPULATIONS

h2o.cut: Convert H2O Numeric Data to Factor.

CHARACTER COLUMN MANIPULATIONS

h2o.strsplit: String Split: "Splits the given factor column on the input split".

h2o.tolower: Convert the characters of a character vector to lower case.

h2o.toupper: Convert the characters of a character vector to upper case.

h2o.trim: Trim spaces: "Remove leading and trailing white space".

h2o.gsub: Match a pattern & replace *all* instances (occurrences) of the matched pattern with the replacement string globally.

h2o.sub: Match a pattern & replace the *first* instance (occurrence) of the matched pattern with the replacement string.

FACTOR LEVEL MANIPULATIONS

h2o.levels: Display a list of the unique values found in a column of categorical data.

DATE MANIPULATIONS

h2o.month: Convert the entries of a H2OParsedData object from milliseconds to months (on a 0 to 11 scale).

h2o.year: Convert the entries of a H2OParsedData object from milliseconds to years, indexed starting from 1900.

MATRIX OPERATIONS

%*%: Multiply two conformable matrices.

t: Given a matrix or data.frame x, t returns the transpose of x.

Data Modeling

MODEL TRAINING: SUPERVISED LEARNING

h2o.deeplearning: Deep Learning Neural Networks.

h2o.gbm: Gradient Boosted Classification Trees and Gradient Boosted Regression Trees.

h2o.glm: Generalized Linear Model, fit by specifying a response variable, a set of predictors, and a description of the error distribution.

h2o.naiveBayes: Naive Bayes Classifier.

h2o.randomForest: Random Forest Classification.

h2o.xgboost: Extreme Gradient Boosted Model.

MODEL TRAINING: UNSUPERVISED LEARNING

h2o.prcomp: Principal Components Analysis.

h2o.kmeans: k-means Clustering.

h2o.anomaly: Detect anomalies using a H2O deep learning model with auto-encoding.

h2o.deepfeatures: Extract the non-linear features using a H2O deep learning model.

GRID SEARCH

h2o.grid: Efficient method to build multiple models with different hyperparameters.

MODEL SCORING

h2o.predict: Obtain predictions from various fitted H2O model objects.

MODEL METRICS

h2o.model metrics: Given predicted values (target for regression, class-1 probabilities, or binomial or per-class probabilities for multinomial), compute a model metrics object.

REGRESSION MODEL HELPER

h2o.mse: Display the mean squared error calculated from a column of predicted responses and a column of actual (reference) responses.

CLASSIFICATION MODEL HELPERS

h2o.accuracy: Between cluster sum of squares.

h2o.auc: AUC (area under ROC curve).

h2o.confusionMatrix: Display prediction errors for classification data (predicted vs reference).

h2o.hit_ratio_table: Retrieve the Hit Ratios.

h2o.performance: Evaluate the predictive performance of a model via various measures.

CLUSTERING MODEL HELPER

h2o.betweeness: Between Cluster Sum of Squares.

h2o.centers: Retrieve the Model Centers.

Cluster Operations

H2O KEY VALUE STORE ACCESS

h2o.assign: Assign H2O hex.keys to R objects.

h2o.getFrame: Get H2O dataset Reference.

h2o.getModel: Get H2O model reference.

h2o.ls: Display a list of object keys in the running instance of H2O.

h2o.rm: Remove H2O objects from the server where the instance of H2O is running, but does not remove it from the R environment.

H2O OBJECT SERIALIZATION

h2o.loadModel: Load H2OModel from disk.

h2o.saveModel: Save H2OModel object to disk.

H2O CLUSTER CONNECTION

h2o.init (nthreads = -1): Connect to a running H2O instance using all CPUs on the host.

h2o.shutdown: Shut down the specified H2O instance. All data on the server will be lost!

H2O LOAD BALANCING

h2o.rebalance: Rebalance (repartition) an existing H2O dataset into given number of chunks (per Vec), for load-balancing across multiple threads or nodes.

H2O CLUSTER INFORMATION

h2o.clusterInfo: Display the name, version, uptime, total nodes, total memory, total cores and health of a cluster running H2O.

h2o.clusterStatus: Retrieve information on the status of the cluster running H2O.

H2O LOGGING

h2o.clearLog: Clear all H2O R command and error response logs from the local disk.

h2o.downloadAllLogs: Download all H2O log files to the local disk.

h2o.logAndEcho: Write a message to the H2O Java log file and echo it back.

h2o.openLog: Open existing logs of H2O R POST commands and error responses on the local disk.

h2o.getLogPath: Get the file path for the H2O R command and error response logs.

h2o.startLogging: Begin logging H2O R POST commands and error responses.

h2o.stopLogging: Stop logging H2O R POST commands and error responses.